

# CS145: Data Management and Data Systems

Stanford University, Fall 2019

Project 3: Querying, Visualizing, Predicting -- The Full Data Cycle  
15% of Course Grade

**Proposal Due Date: Friday, November 8th, 11:59PM**

**Project Due Date: Monday, December 2nd, 11:59PM**

## Overview

Welcome to the final CS145 project! In the first project you learned how to navigate a reasonably complicated dataset and extract information from it. In the second project, you learned how to visualize and reason about the information in these datasets.

In this assignment, you will use the tools you've learned throughout the quarter to follow your own explorations on a topic of your choice. You will pick your own dataset and come up with one or more interesting questions that you want to explore, just like how you began to explore what makes a Git repository popular using the GitHub dataset for Project 2. You will finish by using machine learning to make predictions related to your question.

This project may be done alone or in pairs. However, late days are applied individually. If a pair submits late, any member without late days will receive a zero.

**Note:** We've received feedback that Colaboratory is not a very good interface for collaborative work because it lacks live editing, meaning that partners may accidentally override each other's work. **If you are working in a pair, we recommend either pair programming or having each person work on a separate copy of the notebook and then merging at the very end.** Pairs should still only turn in one notebook.

For this project, we will also require submission of a very short proposal (see **Task A** for more details), detailing your choice of dataset and question as well as your group. The intent of the proposal is for CAs to give feedback on whether or not your dataset and question is appropriate for this project; it is worth 5% of the project grade and must be submitted by **Friday, November 8th** at 11:59 PM. **No late days can be used for this proposal.**

## Task A: Project Proposal (5%)

This proposal is not meant to be a long assignment; you should not need to write much more than two or three small paragraphs. To submit your proposal, go to the *Project 3 - Proposal* assignment and fill out the questions directly on Gradescope, like with Homework 1. If working in pairs, **you should only make one submission for the pair** and then add both group members to the submission. The proposal will ask about what dataset you plan to use and what question(s) you'd like to tackle regarding that dataset.

The proposal is worth 5% of your project grade and is **due on Friday, November 8th**. Don't worry too much about it; as long as you turn it in and fill out the form with reasonable detail, you will receive full credit. After you turn it in, the CA's will quickly go through the proposals to ensure that the dataset and question you choose to explore are rich and complex enough for this project; by Tuesday, November 12th, you will have feedback on your proposal. We expect that in most cases, it will be a simple go-ahead; for some, you may be asked to choose a dataset with greater size or complexity or to choose a question with more depth. In other cases, we may advise you to decrease the scope of your project.

You are allowed to change your choice of dataset and question after you submit your proposal; however, you will have to take initiative and come to OH if you want to verify that your new choices are appropriate for this project.

## Task B: Introduction to Machine Learning (10%)

For a small portion of this project, you will be using BigQuery's machine learning features to make predictions related to your question of choice. Do not worry if you've never studied machine learning before! We have created a Colab notebook, `project3_ml_warmup.ipynb`, that will guide you through the fundamentals of what you need to know to complete this project. You can access the notebook from the course website, **make a copy of it in your own drive**, and begin the assignment. Read the notebook and fill out the questions where specified.

This notebook will be worth 10% of your project grade.

## Task C: Your Own Data Cycle (85%)

For the bulk of Project 3, you will be exploring a question that interests you on a BigQuery dataset of your choice. Please create a Colab notebook containing all of your work. You will use what you've learned from Project 1 and Project 2 to begin your explorations on the dataset and your question. Once you've done your explorations, you will use machine learning to make predictions relevant to your interesting question.

As an example, take the GitHub dataset from Project 2. We were trying to answer the question of what factors impact the popularity of a GitHub repository. Given this question, once we finish

our explorations and decide what features were important, we can use machine learning to predict the popularity of various repositories based on those features.

This part of the project is largely open-ended; the specific explorations you choose to do are up to you to decide. However, we expect your project will contain at least the following sections:

- **Project overview (5%)**
  - Provide a brief project overview; that is, an explanation of the main question you are trying to answer.
  - Briefly summarize any supplementary questions that will be used to help answer central questions.
- **Analysis of your dataset (10%)**
  - Provide a brief overview of what your dataset is about as well as the **overall size and complexity** (in terms of bytes, the number of tables, etc.). Keep in mind that small datasets may lead to difficulty when attempting machine learning.
  - Comment on how the dataset you chose is organized. Some questions you may think about are: **At a high level, what information does each table capture? What are the relationships between tables / how are they related?** You may want to elaborate by providing keys and foreign keys of tables as necessary.
  - Your analysis should demonstrate that you understand the structure and the content of the data you are working with. You may supplement your analysis with concrete examples to corroborate your statements.
- **Exploring your questions, with appropriate visualizations (55%)**
  - Write queries to gather information about your dataset -- you may use any plotting library of your choice to visualize and understand your data. Ask and answer quantitative questions that revolve around your central questions.
  - Your exploration should both shed light on the questions that you've asked yourself and provide insights for the feature engineering you will use when you generate predictions.
  - In this section, conclude with analysis and summary of your observations. What features of your data seem especially prominent or related to your question? Are there anomalies? What trends do you see? Have you answered your questions?
  - Your analysis in this section should demonstrate that you have a reasonable quantitative understanding of your dataset. By the end of this section, you should be ready to use your newly-gained domain knowledge to generate predictions.
- **Predictions based on your explorations (20%)**
  - Based on your explorations and analysis from the previous section, train a prediction model using BigQuery for the questions you would like to answer.
  - Evaluate your model in BigQuery. Comment on the performance of your model. Is it good? Is it bad? In either case, why do you think your model does well or badly? What metrics are you using to measure performance and why?
  - Finally, use your model to make predictions **on data besides that which you trained (i.e. created your model) on.**
- **Conclusion (10%)**

- What have you learned? What conclusions have you made or been unable to make about your dataset and why? What is obvious, and what did you not expect to see? If you had more time, what other data exploration would you pursue?

A more detailed rubric is available on the course website, along with a very simple template notebook. You may choose to use and modify this template as you please.

Note that we are generally more concerned with the depth with which you've leveraged visualizations and BigQuery to gain insight into your questions than the quality of your machine learning predictions. This is an open-ended project -- there is no right or wrong answer, only quality of exploration and inquiry.

This part will be worth 85% of the project grade.

## Honor Code

As in all Stanford classes, you are expected to follow the Stanford Honor Code. For example, the following activities are prohibited and will be treated as Honor Code violations (this is not intended to be a complete list of Honor Code violations):

- Submitting code that you did not write personally, with the exception of project code written by your partner.
- Consulting pre-existing solutions for problem sets and projects (such as solutions posted on the Internet).
- Posting your solutions on the Internet or making them available to other students in any form.

You are allowed to discuss general approaches and issues with other students in the class besides your project partner. It's also fine to give other students help finding bugs if they are stuck, or to answer general questions. But, any code you write must be written by you and your partner, from scratch, without consulting existing solutions. We reserve the right to use computer software such as MOSS to analyze material that you submit in order to detect duplication with other students or existing solutions.

A general way to think about this is that if a particular activity significantly short-circuits the learning process (it saves you time but reduces the amount you learn and/or figure out on your own), or if it misrepresents your capabilities or accomplishments, then it is probably an Honor Code violation.

## Submission Instructions

### Proposal Submission:

Please directly answer the proposal questions in the *Project 3 - Proposal* assignment on Gradescope.

### Project Submission:

1. Download the Machine Learning Warmup Colab notebook as an iPython notebook - you can do this by going to *File > Download .ipynb*. Submit it to the *Project 3 - ML Warmup iPython* assignment on Gradescope.
2. Create a PDF of the Machine Learning Warmup Colab notebook, **making sure that you have run all cells first**. Make sure you've closed the table of contents sidebar before you create the PDF so we can easily see your work and output. Submit this PDF to the *Project 3 - ML Warmup PDF* assignment on Gradescope.
3. Download your personal exploration Colab notebook as an iPython notebook. Submit the iPython notebook to the *Project 3 - Personal Explorations iPython* assignment on Gradescope. Please name your file according to the format `firstname_lastname.ipynb` for single-person submissions, and for two-person submissions `firstname1_lastname1_firstname2_lastname2.ipynb` (e.g. `jennie_chen.ipynb` or `jennie_chen_qiwen_wang.ipynb`). There is no need to submit a PDF of this notebook.
4. For each of your submissions on Gradescope, **make sure to add your project partner as a group**. You can do this after submitting by clicking *View or Edit Group* underneath your name and searching for your partner.

In total, for this project (ignoring the proposal) you should be submitting to one file to each of three different Gradescope assignments.

**Note:** We reserve the right to deduct points from your project if you do not follow the submission instructions. Please leave yourself enough time to do the submission!

You may resubmit as many times as you like; however, only the latest submission and timestamp will be saved, and we will use your latest submission for grading your work and determining any late penalties that may apply. Submissions via email will not be accepted.